

Forecasting Daily Online Retails Sales using SARIMAX model

2025-02-04

Table of Contents

1 Main Objective of the Analysis	1
2 Brief Description of the Dataset	2
3 Data Exploration and Feature Engineering	2
3.1 Data Exploration	2
3.2 Feature Engineering	2
4 Data Anlaysis	3
5 Methodology	4
6 Results	5
7 Conclusion	6
8 Key Findings and Insights	6
9 Suggestions for Next Steps	7

1 Main Objective of the Analysis

The main objective of this analysis is to forecast future sales of retail products based on historical sales data. The time series forecasting model selected is a SARIMA (Seasonal AutoRegressive Integrated Moving Average) model, which is well-suited for time series data that exhibits both trend and seasonality, such as retail sales data. By using this model, the business can gain valuable insights into future sales, enabling better inventory management, demand planning, and resource allocation. The analysis aims to help stakeholders, such as supply chain managers, inventory planners, and marketing teams, make informed decisions based on forecasted sales trends.

2 Brief Description of the Dataset

The dataset chosen for this analysis is the Online Retail dataset from the UCI Machine Learning Repository. The dataset contains transactional data for a UK-based retail company, detailing purchases made by customers between 2010 and 2011. The key attributes in the dataset include:

- InvoiceNo: Unique identifier for each invoice.
- StockCode: Product code for each item purchased.
- Description: Product description.
- Quantity: Number of units purchased.
- InvoiceDate: Date and time when the invoice was generated.
- UnitPrice: Price per unit of the product.
- CustomerID: Unique identifier for customers.
- Country: Country of the customer.

The primary objective of this analysis is to forecast the total sales ($\text{Quantity} * \text{UnitPrice}$) for the company over a specified future period. This forecast will aid the company in predicting demand and adjusting its strategies for inventory and marketing.

3 Data Exploration and Feature Engineering

3.1 Data Exploration

The data exploration process involved inspecting the dataset for any obvious issues or inconsistencies

There were no missing values or outliers. So we proceed with the analysis

3.2 Feature Engineering

- The InvoiceDate column was parsed into datetime format to facilitate time series analysis.
- The time series data was aggregated on a daily basis to calculate total sales for each day, considering the product of Quantity and UnitPrice and Sales column was created.

4 Data Anlaysia



Figure 1: Series, Trend, Seasonality, Residuals Diagram

Figure 1 shows the decomposition of the components of the Time Series. We can observe a slight trend and a clear weekly seasonal pattern.




Figure 2: ACF and PACF plots

Figure 2 show the ACF and PACF plots. Here also we can clearly observe that there is a weakly seasonality, as a result of seasonal lags been significant.

5 Methodology

The forecasting model for the daily sales data was built using two distinct approaches: SARIMAX and Stepwise ARIMA. Both models were aimed at capturing the temporal dependencies and seasonality inherent in the data to make accurate predictions. The steps followed for each method are outlined below:

1. SARIMAX Model:

- **Model Definition:** A SARIMAX model was specified with the following parameters:
 - **AR(1):** One autoregressive term to capture the dependence of the current sales on the previous day's sales.
 - **MA(1) Seasonal Component:** A seasonal moving average term at lag 7 to account for weekly seasonal patterns in the sales data.
 - **Trend:** A constant trend was added to the model to capture any potential upward or downward trends over time.
- **Model Fitting:** The SARIMAX model was fit using the `sm.tsa.statespace.SARIMAX` function from the `statsmodels` library. The model was trained on the historical data, and the optimal parameters were estimated using maximum likelihood estimation.

2. Stepwise ARIMA Model:

- **Model Definition:** The stepwise ARIMA model was chosen using the `auto_arima` function from the `pmdarima` library. This function automates the process of selecting the best ARIMA model by evaluating different combinations of parameters:
 - **AR(p), MA(q):** A range of values from 1 to 3 were tested for both AR and MA components.
 - **Seasonal Parameters (P, Q):** The model was defined to have a seasonal component with a period of 7 (weekly seasonality), with seasonal MA terms considered at lags 7 and 14.
 - **Differencing (d, D):** The model applied a non-seasonal differencing of order 0 and seasonal differencing of order 1 to make the series stationary.
- **Model Fitting:** The `auto_arima` function iterated through all possible combinations of parameters and selected the best model based on the AIC score. This resulted in the optimal configuration of ARIMA(1,0,1)(0,1,2)[7].

3. Model Evaluation:

- Both models were evaluated based on their **AIC (Akaike Information Criterion)** and **BIC (Bayesian Information Criterion)** values, with lower values indicating better model fit. Additionally, the **Ljung-Box test** was applied to the residuals of both models to check for autocorrelation and ensure that the models had adequately captured the temporal dependencies in the data.

By following these steps, the two models were built and evaluated to identify the best approach for forecasting daily sales. The SARIMAX model was initially fit to the data, but the stepwise ARIMA model, with its automatic parameter selection, provided a more accurate and parsimonious model.

6 Results

1. SARIMAX Model:

```
Dep. Variable:  Sale      No. Observations:   374
Model:  SARIMAX(1, 0, 0)x(0, 1, [1], 7) Log Likelihood  -3994.175
Date:    Tue, 04 Feb 2025      AIC 7996.351
Time:    14:20:07      BIC 8011.972
Sample:  12-01-2010  HQIC      8002.557
- 12-09-2011
Covariance Type:  opg
coef    std err  z    P>|z|    [0.025  0.975]
intercept    119.1026    237.682  0.501    0.616    -346.745    584.951
ar.L1    0.2701  0.069    3.899    0.000    0.134    0.406
ma.S.L7  -0.7873  0.048   -16.521  0.000   -0.881   -0.694
sigma2  2.215e+08    0.002   1.42e+11  0.000   2.22e+08    2.22e+08
Ljung-Box (L1) (Q):  0.74      Jarque-Bera (JB):   328.25
Prob(Q):    0.39      Prob(JB):    0.00
Heteroskedasticity (H): 1.80      Skew:    1.07
Prob(H) (two-sided):    0.00      Kurtosis:    7.11
```

The SARIMAX model, with one autoregressive (AR) term, one seasonal moving average term (MA) at lag 7, and a constant trend, produced an AIC of 7996.351 and a BIC of 8011.972.

The Ljung-Box test revealed no significant autocorrelation in the residuals (p-value = 0.39), indicating that the model effectively captured the patterns in the data. Despite the absence of residual autocorrelation, the higher AIC and BIC values compared to the stepwise ARIMA model suggest that the model's performance can be improved

2. Stepwise ARIMA Model:

```
Dep. Variable:  y    No. Observations:   374
Model:  SARIMAX(1, 0, 1)x(0, 1, [1, 2], 7)  Log Likelihood  -3984.125
Date:    Tue, 04 Feb 2025    AIC 7980.251
Time:    16:11:56    BIC 8003.683
Sample:  12-01-2010  HQIC    7989.561
- 12-09-2011
Covariance Type:  opg
coef    std err z    P>|z|    [0.025  0.975]
intercept    56.6926  72.460  0.782    0.434   -85.327  198.713
ar.L1    0.8607  0.113  7.598  0.000    0.639  1.083
ma.L1   -0.6928  0.152  -4.557  0.000   -0.991  -0.395
ma.S.L7 -0.7252  0.076  -9.496  0.000   -0.875  -0.576
ma.S.L14 -0.1233  0.074  -1.657  0.098   -0.269  0.023
sigma2  2.213e+08  8.55e-05  2.59e+12  0.000  2.21e+08  2.21e+08
Ljung-Box (L1) (Q): 0.08    Jarque-Bera (JB): 379.84
Prob(Q): 0.77    Prob(JB): 0.00
Heteroskedasticity (H): 2.04    Skew: 1.17
Prob(H) (two-sided): 0.00    Kurtosis: 7.40
```

The stepwise ARIMA model identified a combination of one AR term, one MA term, and two seasonal MA terms (at lags 7 and 14) as the optimal configuration. This model achieved a lower AIC of 7980.251 and a BIC of 8003.683, signifying a better fit to the data compared to the SARIMAX model. The Ljung-Box Q-test (p-value = 0.77) confirmed that the residuals were white noise, further supporting the adequacy of the model. The inclusion of the seasonal MA terms (lag 7 and lag 14) helped capture additional seasonality that the SARIMAX model could not

7 Conclusion

In conclusion, the stepwise ARIMA model outperforms the SARIMAX model (1, 0, 1)(0, 1, 2)[7] based on the AIC, BIC, and residual diagnostics, making it the preferred model for forecasting future sales.

8 Key Findings and Insights

- **Seasonality:** The dataset showed strong weekly seasonality, which was effectively captured by the SARIMA model. This suggests that certain days of the week consistently perform better in terms of sales.
- **Trend:** The sales data exhibited a positive trend, indicating growing demand over time.

- **Forecasting Future Sales:** The SARIMA model provided reliable forecasts of future sales, with the ability to estimate the impact of seasonal patterns. These findings provide the business with actionable insights for resource planning, inventory management, and marketing strategies.

9 Suggestions for Next Steps

Incorporating Deep Learning: Deep learning models (e.g., LSTM or RNN) could provide better result